

# Large-scale computing resources at CU Boulder

Jonathon Anderson and Shelley Knuth

[rc-help@colorado.edu](mailto:rc-help@colorado.edu)

[www.colorado.edu/rc](http://www.colorado.edu/rc)

Slides:

[https://github.com/ResearchComputing/CEASE\\_Webinar\\_2021](https://github.com/ResearchComputing/CEASE_Webinar_2021)

# Outline

- Research Computing
  - Resources
  - User facing services and support
- CRDDS
  - Services and support
- Rocky Mountain Advanced Computing Consortium
  - Leveraging this collaboration
- Procurement of our new computing system

# What is Research Computing?

- Provide services for researchers that include:
  - Large scale computing
  - Data storage
  - High speed data transfer
  - Data management support
  - Consulting
  - Training
- We are likely best known for:
  - Summit Supercomputer
  - PetaLibrary storage

# What Would I Use Summit For?

- Research Computing is more than just Summit
- What would you use Summit For?
  - Solving large problems that require more:
    - Memory than you have on your personal computer
    - Cores/nodes/power than you have on your personal computer
  - High performance GPU computing
  - High memory jobs
  - Visualization rendering
- Not a place for:
  - Large data storage

# RMACC Summit

- 67% CU, 23% CSU, 10% RMACC
- 500+ compute nodes
  - Majority 24 cores / node “Haswell” (shas) nodes
- 13,000+ cpu cores
- Intel Omni-Path Architecture interconnect
- GPFS scratch file system
  - 1.2 PB




# RMACC Summit node types

- 10 Nvidia K80 GPU nodes (sgpu)
  - 2 Nvidia Tesla K80 accelerators / node
- 5 “high-memory” nodes (smem)
  - 2 TB of memory / node
  - 48 cores / node
- 20 Intel Xeon “Knights Landing” Phi nodes (sknl)
  - 68 cores / node

# Blanca condo compute cluster

- Aggregate cluster from researcher capital purchase contributions
- Pervasively heterogeneous
- 200+ nodes and almost 3,400 CPU cores
- 30 distinct contributors
- Oldest node 2012; newest node 2021



# PetaLibrary


- Active storage on ZFS (hard disk)
  - \$45/TB/year
  - Accessible from all RC compute nodes
- Archive storage on LTFS (tape)
  - \$20/TB/year
  - iRODS





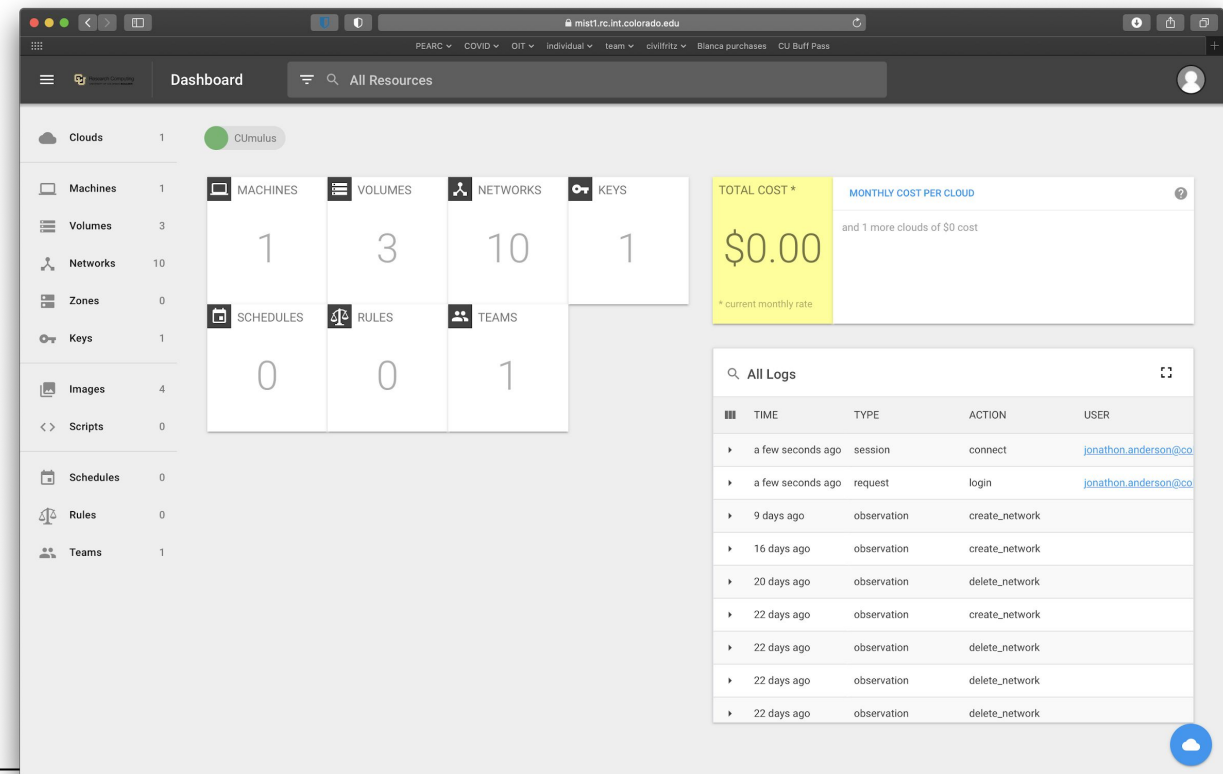
# EnginFrame

- 3D accelerated virtual desktop
- Web-accessible
- Backed by 2 visualization nodes
  - 2 Nvidia Quadro RTX 8000 accelerators / node
  - 48 GiB memory / accelerator



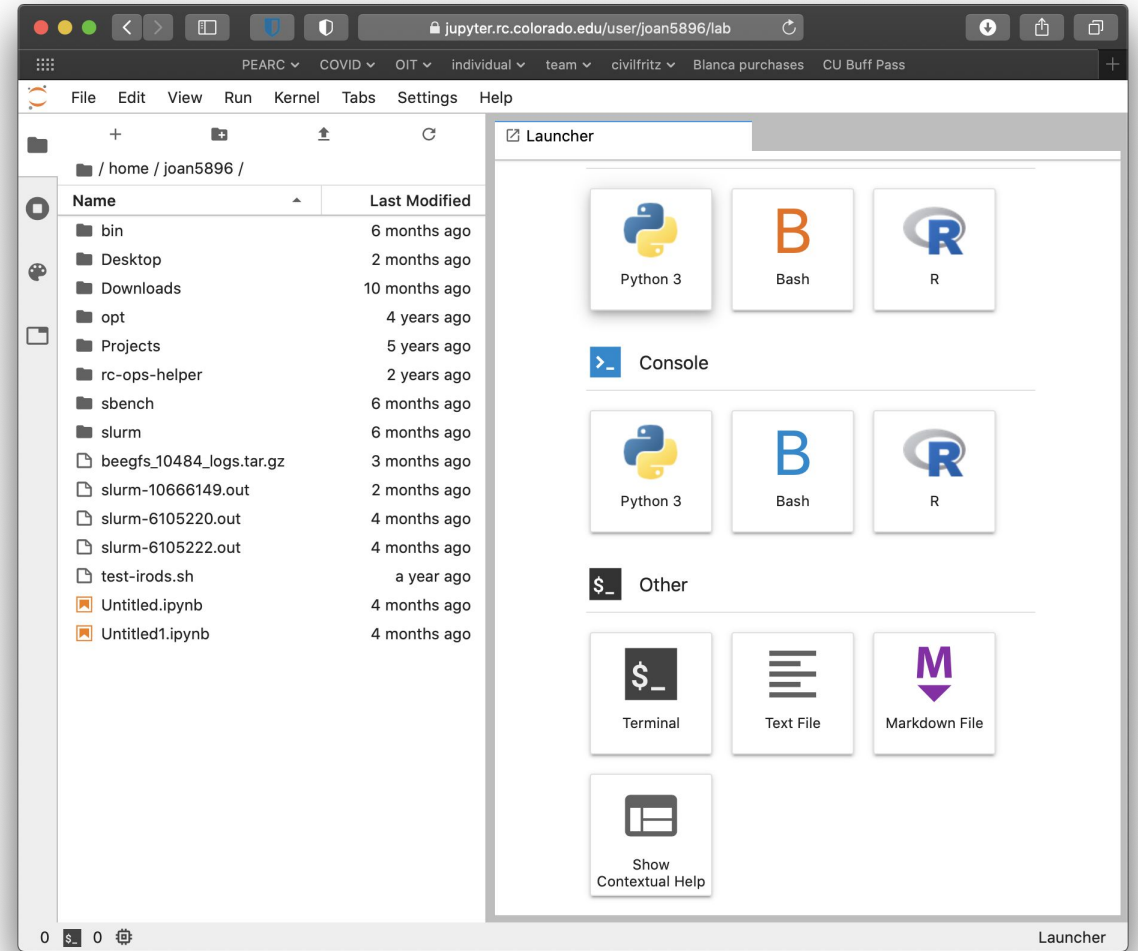
# CUmulus Hybrid Cloud

- On-premises OpenStack “IaaS” cloud
- Access to commercial clouds
- mist.io multi-cloud interface
- Hosting for
  - Database servers
  - Science Gateways
  - Compute pipelines



# JupyterHub

- Python, R/RStudio, Bash
- Supports user-provided custom kernels
- Able to run on RMACC Summit and Blanca



# **“CURC-3” next-generation HPC**

- Designed for regular expansion
- Estimated compute:
  - 160 CPU compute nodes (7,680 cores)
  - 20 GPU compute nodes  
(60 accelerators, mixture of Nvidia and AMD)
- Anticipated partitioned HDR-100 interconnect
- Anticipated 25Gb Ethernet to storage (including PetaLibrary)
- GPFS scratch
  - 22 GB/s
  - 2 PB

# Faculty Support

- Ease access to resources
  - User support services
- CRDDS
  - Review of data management plans for external grants
  - Data storage and publication
    - CU Scholar
    - PetaLibrary
- Facilities statements
- Staff support for grant work



# Description of User Support Services

- Trainings
- Consultations
- Ticketing system
- Office hours
- Outreach
- Group configuration
  - Team lead, three full time staff, one 20% staff member
- Approximately 800 active users





# Trainings

- 2019-2020 offered nearly 200 trainings
  - "Introduction to Python"
  - "Containerization for Research and Development Applications"
  - "Fundamentals of HPC"
- Over 1000 attendees
- Online option
- Topics chosen by request or by staff



# Office Hours/Consulting

- Began in mid-2018
- Collaboration with CRDDS
- Tuesdays 12-1 pm, Thursdays 1-2 pm
  - RC only available on Tuesdays
- 400 consultations per year
- Fully online right now
- Normally in person
- 1:1 consultations also available





Supporting interdisciplinary data-intensive research & education

= CU Center for Research Data and Digital Scholarship (CRDDS)

<https://www.colorado.edu/crdds/>

[crdds@colorado.edu](mailto:crdds@colorado.edu)

- Partnership between University Libraries and Research Computing
- Partnerships with other entities that have a similar mission
- Support of faculty & students with data related topics
- Center space (Norlin E206)


# Data management



- Curation and DOI registration for 219 data sets in the CU Scholar institutional repository in first year
- DataCite DOI agreement with National Solar Observatory (NSO) and Laboratory for Atmospheric and Space Physics (LASP)
- CoreTrustSeal repository certification in collaboration with National Center for Atmospheric Research (NCAR)
- Data Publishing of large data sets
  - Metadata in CU Scholar linked through Globus to data sets on PetaLibrary
- Support of grant proposals

# Data Storage and Publishing

- Meet funding agency or publication demands
- PetaLibrary
  - Large-scale
  - No sensitive data
- CU Scholar
  - Under 10 GB
  - DOI
  - Freely accessible
- Combination
  - PetaLibrary backend
  - Publish larger datasets



# RMACC

the Rocky Mountain Advanced Computing Consortium

- Group of institutions in Rocky Mountain west
- NSF projects
  - MRI grant to procure Summit
  - Cyberteam project
    - Data workflows
    - Access to HIPAA compliant data storage
  - Hybrid cloud
- Share of Summit – 10%
- Share of hybrid cloud - 20%
- 70 users across six institutions
- Ongoing efforts to reach out
- CU provides user support – not dedicated
- Funding agencies love collaborations!



# Questions?

Email:

Research Computing [rc-help@colorado.edu](mailto:rc-help@colorado.edu)

CRDDS [crdds@colorado.edu](mailto:crdds@colorado.edu)

Website: [www.colorado.edu/rc](http://www.colorado.edu/rc)

Slides:

[https://github.com/ResearchComputing/CEASE\\_Webinar\\_2021](https://github.com/ResearchComputing/CEASE_Webinar_2021)